

# POV: Identify Grid stability risks in advance using Notebooks in MS Fabric

Prepared by LSHCERU BG  
Muskan Singhal



# Objective

- ✓ To monitor and predict grid stress levels based on real-time and forecasted weather conditions, enabling proactive load balancing and outage prevention.
- ✓ To integrate weather forecasts into grid analytics for dynamic risk scoring and early warning of instability events.
- ✓ To develop predictive models that correlate weather anomalies (e.g., heatwaves, wind speed, storms) with grid performance metrics across regions.
- ✓ To visualize the impact of weather variables on grid reliability through interactive dashboards and alert systems.



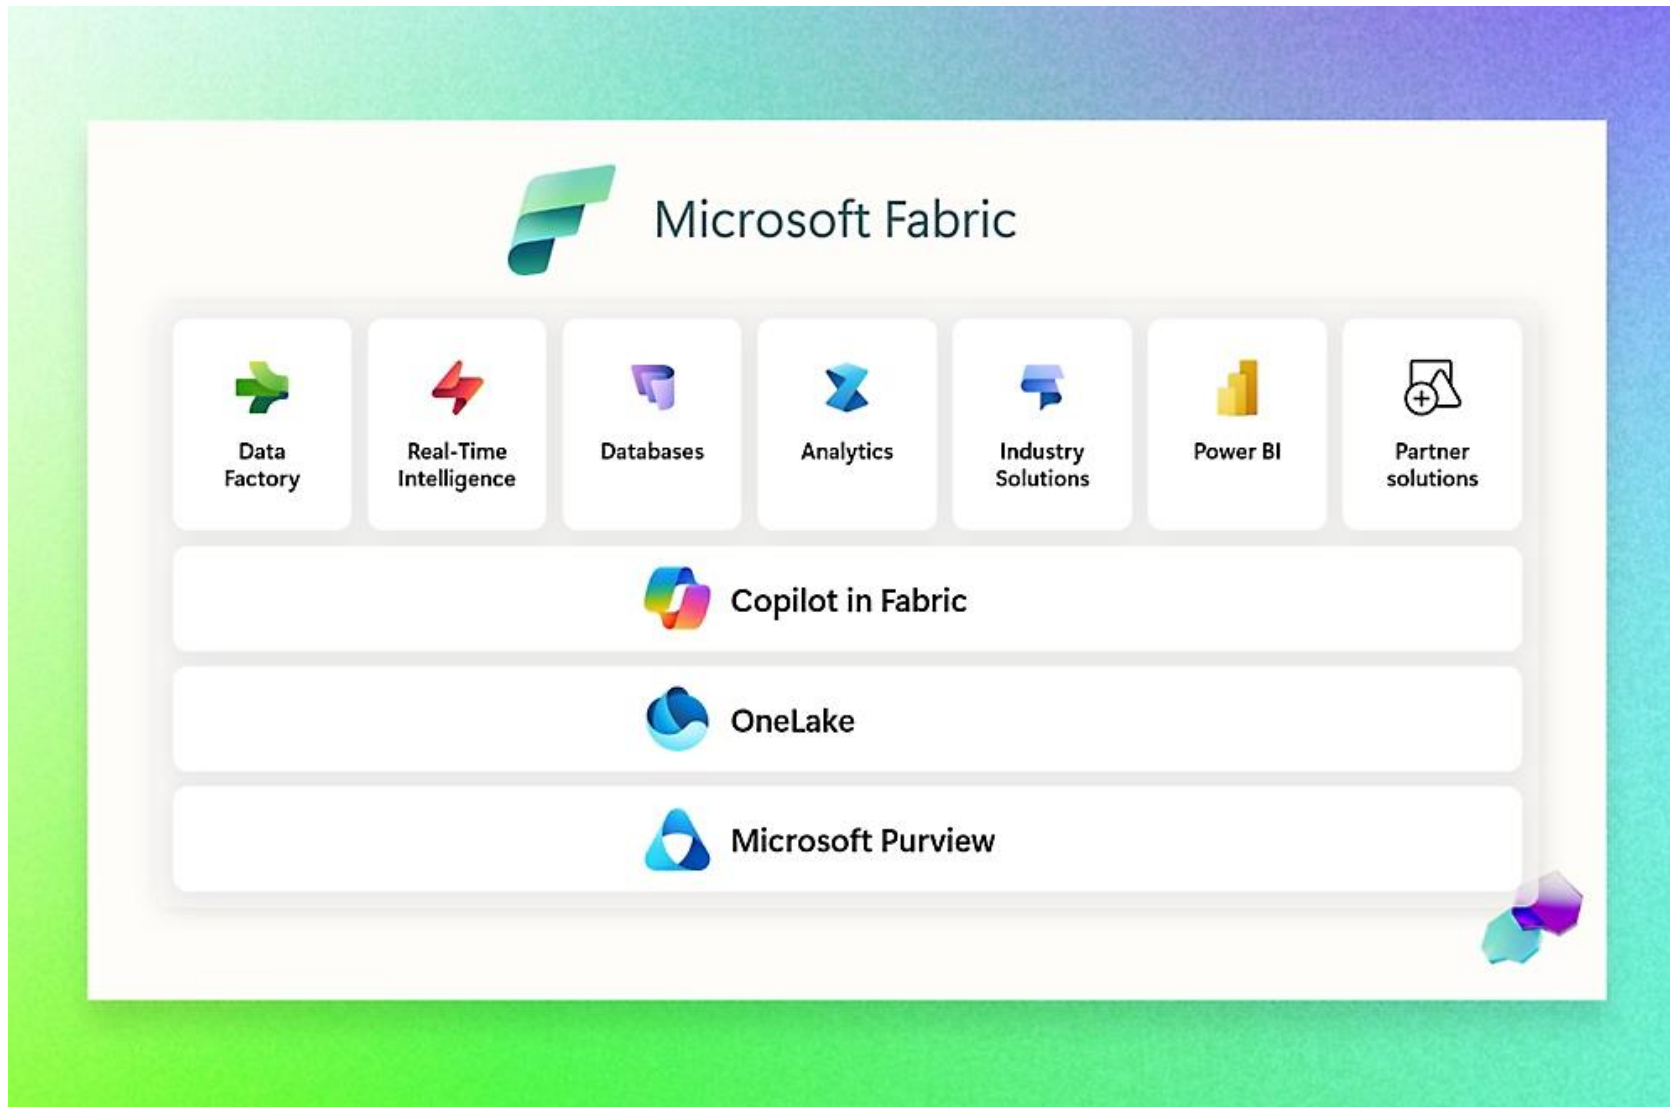
# Overview of Grid stability risks

It's about using **real-time and forecasted weather data** to anticipate conditions that could destabilize the power grid – before they happen. This enables utilities and grid operators to take **preventive actions** like load balancing, rerouting, or activating backup systems.

## Weather Factors That Impact Grid Stability

Weather Element	Risk to Grid Infrastructure
High Temperature	Overheating of transformers and substations
Strong Winds	Damage to transmission lines and poles
Heavy Rain/Floods	Substation flooding, short circuits
Snow/Ice	Line sagging, equipment failure
Lightning	Surge damage, outages
Solar Radiation	Impacts solar generation forecasting

# Microsoft Fabric | Ecosystem



# What is a Notebook in Microsoft Fabric?

The notebook serves as an interactive coding environment where users can use languages like Python (PySpark, Pandas) to perform tasks such as data ingestion, transformation, exploration, and analysis by loading data into dataframes for processing.

- A web-based interactive coding environment for data engineering and data science.
- Supports multiple languages: Python, R, Scala, and SQL.
- Ideal for data ingestion, transformation, visualization, and machine learning.
- Integrated with Lakehouse, pipelines, and other Fabric experiences.
- Offers rich visualizations, markdown support, and enterprise-grade security.

## With a Fabric notebook, you can:

- Get started with zero set-up effort.
- Easily explore and process data with intuitive low-code experience.
- Keep data secure with built-in enterprise security features.
- Analyze data across raw formats (CSV, txt, JSON, etc.), processed file formats (parquet, Delta Lake, etc.), using powerful Spark capabilities.
- Be productive with enhanced authoring capabilities and built-in data visualization.

# Core Features of Fabric Notebooks

Feature	Description
<b>Multi-language Support</b>	Write code in Python, Spark SQL, and markdown – all in one place.
<b>Apache Spark Integration</b>	Run distributed data processing jobs natively with Spark clusters.
<b>Lakehouse Connectivity</b>	Read/write directly to Lakehouse tables using shortcuts and workspace links.
<b>Markdown Documentation</b>	Add rich text, headers, and formatting to explain logic and annotate steps.
<b>Built-in Visualizations</b>	Use Python libraries (e.g., matplotlib, seaborn) or Spark display functions for charts.
<b>Pipeline Integration</b>	Trigger notebooks as activities in Fabric Data Factory pipelines.
<b>Scheduler Support</b>	Automate notebook runs with time-based triggers and security context control.
<b>Security Context Awareness</b>	Runs under user, pipeline owner, or scheduler identity – critical for governance.
<b>Versioning &amp; Export</b>	Save, export, and track notebook changes for audit and collaboration.
<b>Real-time Collaboration</b>	Co-edit notebooks with team members in shared workspaces.

# ETL Tools used in Microsoft Fabric

## Lakehouse

- Served as the unified repository for raw, cleaned, and enriched weather data across multiple locations and timestamps.
- Structured Lakehouse into staging (raw API ingests) and final zones (transformed, validated datasets) for clear ETL separation.
- Seamlessly integrated with Spark Notebooks for Cross-Tool Compatibility.
- Provided fast, scalable access to weather data for Power BI dashboards and Spark transformations.

## Spark Jobs

- Production grade batch processing jobs that need to run on schedule or be triggered by events.
- Can handle massive datasets & complex transformations using the full spark ecosystem.

## Shortcuts

- Enabled seamless access to curated Lakehouse tables across multiple Fabric workspaces without duplicating data.
- Made curated datasets easily discoverable for Power BI reports without exposing raw ingestion zones.
- Enables cross-workspace access to curated datasets for scalable reuse and governance.

## Notebooks

- Using PySpark, Scala , R, or SQL within Fabric Notebooks
- Gives complete programmatic control over Transformations.
- Ideal for complex business logic, machine learning preprocessing
- Exploit full power of distributed computing.
- Performs timestamp conversion, data cleaning, and enrichment across multi-location feeds.

## Delta Tables

- Stores versioned, query-optimized weather data for downstream analytics and dashboards.

## Real-time Analytics(KQL)

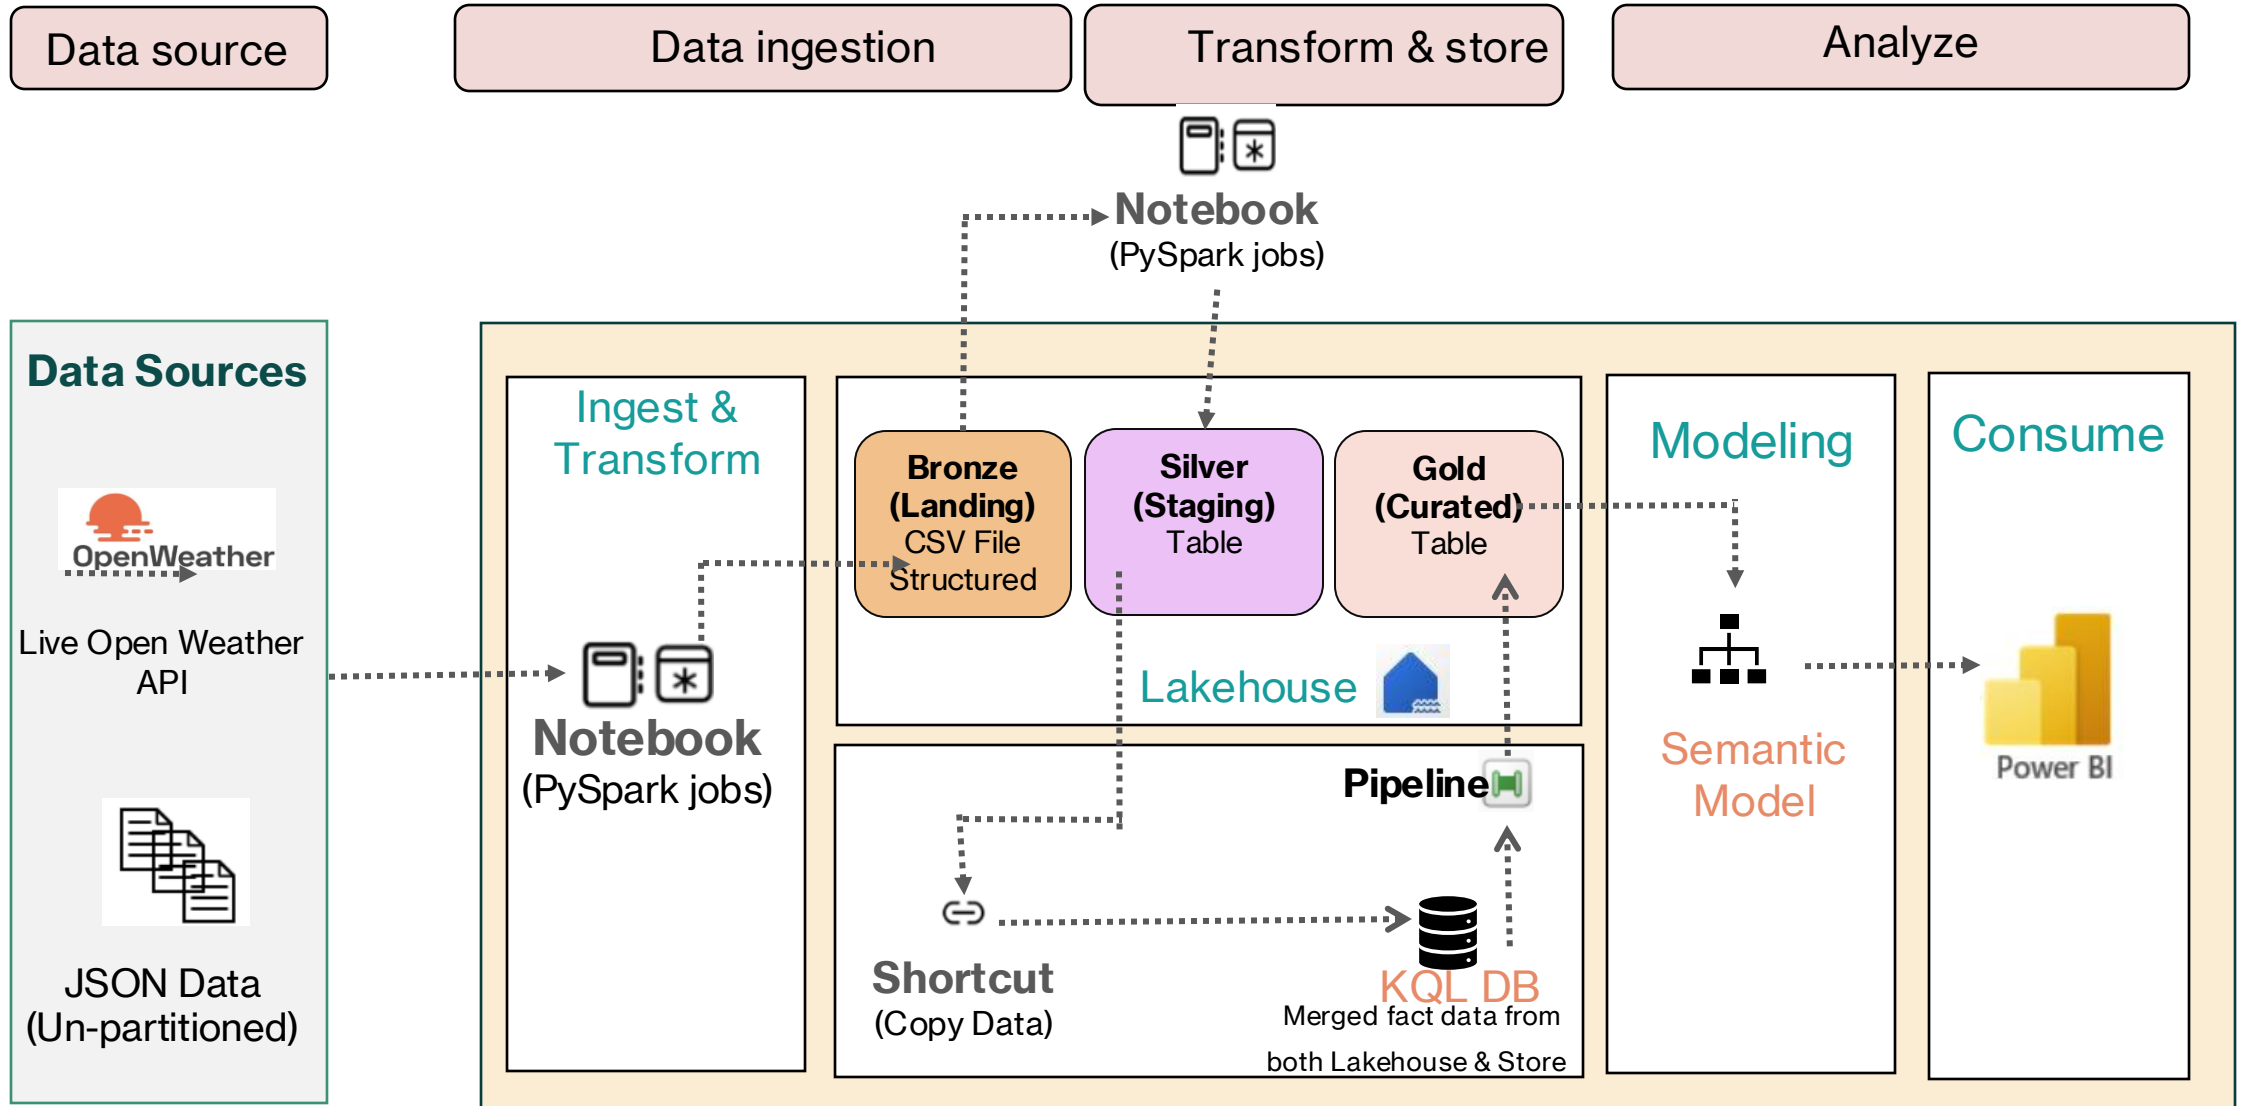
- For streaming data scenarios which need continuous ingestion & transformation.
- KQL provides powerful real-time processing capabilities

## Pipeline

- **Source Connection:** Connects to diverse data sources (e.g., REST APIs, Azure Blob, SQL DBs, OneLake shortcuts).
- **Activities:** Includes Copy Data, Data Flow, Notebook, Spark Job, and custom scripts.
- **Triggers:** Supports manual, scheduled, and event-based triggers for pipeline execution.
- **Parameters:** Enables dynamic pipeline behavior using runtime parameters for flexibility.



# Architecture in Microsoft Fabric





# Create Notebooks-

## 1. Create a Lakehouse

You create a Lakehouse first to serve as a centralized, structured data storage for notebooks to access & manipulate. Notebooks are then used to query this data in the Lakehouse, performing tasks like data ingestion, transformation, analysis, and developing machine learning models. The Lakehouse provides a unified environment, allowing notebooks to interact directly with data without complex data movement, making data exploration and development efficient.

Workspace-> New Item-> Lakehouse

2. Create a new notebook - You can easily create a new notebook from the Fabric Data Engineering homepage, the workspace New option, under workspace.

Workspace-> New Item-> Notebook

## 3. Connect Lakehouse and notebooks

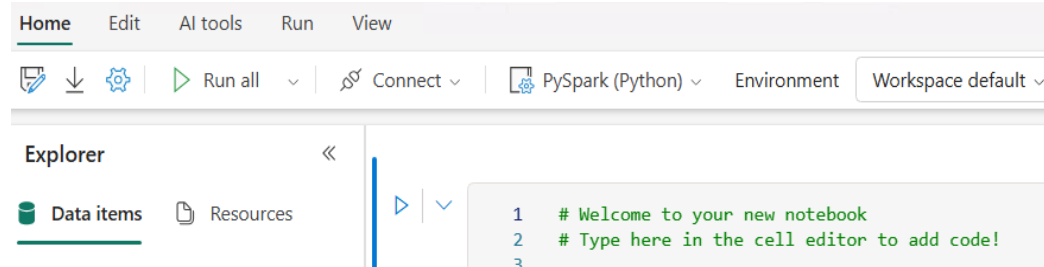

Fabric notebooks support close interactions with Lakehouse; you can easily add a new or existing Lakehouse from the Lakehouse explorer. This connection allow users to build data pipelines, train models, and gain insights from large datasets without needing to import or copy the data. So, We connect a Lakehouse to a notebook to programmatically access and manipulate data stored in the Lakehouse for data engineering, analytics, and AI workloads.

Add Data items-> Existing data sources-> select your Lakehouse-> connect


## 4. Query your PySpark/Spark SQL in Notebook

### Store data

Organize, query, and store your ingested data in an easily retrievable format.



### No data sources added



## How I utilized Notebooks for Weather Data Use Case-

### Step 1. Calling API with Parameters in Notebook using PySpark

In 1st cell of notebook, I making a request to call API by passing API URL with passing few parameters [longitude & latitude] to fetch the data of all available locations available in API Data.

```
Weather at (51.5, -0.09): overcast clouds
Weather at (48.85, 2.34): clear sky
Weather at (28.55, 77.16): haze
Weather at (19.13, 72.89): mist
Weather at (-33.86, 151.2): clear sky
Weather at (13.08, 80.16): scattered clouds
Weather at (12.97, 77.6): broken clouds
```

### Step 2. Checking no. of columns with column names

In 2nd block, I am checking the format of API data whether it is in JSON/Parquet. Also, I used Python Print code to get all the columns and length of columns available in this data.

```
{'coord': {'lon': 77.6, 'lat': 12.97}, 'weather': [{'id': 803, 'main': 'Clouds', 'description':
'broken clouds', 'icon': '04n'}], 'base': 'stations', 'main': {'temp': 298.81, 'feels_like':
299.42, 'temp_min': 298.22, 'temp_max': 299.07, 'pressure': 1009, 'humidity': 76, 'sea_level':
1009, 'grnd_level': 911}, 'visibility': 6000, 'wind': {'speed': 0, 'deg': 0}, 'clouds': {'all':
75}, 'dt': 1758114757, 'sys': {'type': 1, 'id': 9205, 'country': 'IN', 'sunrise': 1758069525,
'sunset': 1758113386}, 'timezone': 19800, 'id': 1277333, 'name': 'Bengaluru', 'cod': 200}

['weather', 'base', 'visibility', 'timezone', 'id', 'name', 'cod', 'coord.lon',
'coord.lat', 'main.temp', 'main.feels_like', 'main.temp_min', 'main.temp_max',
'main.pressure', 'main.humidity', 'main.sea_level', 'main.grnd_level', 'wind.speed',
'wind.deg', 'clouds.all', 'sys.type', 'sys.id', 'sys.country', 'wind.gust',
'Sunrise_UTC', 'Sunset_UTC']
Total columns: 26
```

**Step 3.** Transform the JSON response to a DataFrame using Pandas

In 3rd block, I import Pandas library & created DataFrame because the data we extracted from API is Semi-Structured which we stored in a raw output as a JSON or binary file. So, we can convert the raw output into structured format we can easily store it in Delta Table. Also, I added static column which shows status like Current.

Now the data is shown in a row column format but still not in a fine table.

**Step 4.** Write the Data Frame to a Delta Table

In 4th block, I format the above output in a fine table which shows the data properly.

**Step 5.** Filter the table data

Format the unsupported column names like space, brackets & correct the format of sunrise, sunset date time data which we were getting in Unix Timestamp earlier.

**Step 6.** Export Data Frame into CSV format

**Step 7.** Save that csv file in my Lakehouse

**Step 8.** Loading table into existing Lakehouse

**Step 9.** We can refresh & check in Lakehouse, the table is there.

id	email	first_name	last_name	avatar
7	michael.lawson@re...	Michael	Lawson	https://reqres.in...
8	lindsay.ferguson@...	Lindsay	Ferguson	https://reqres.in...
9	tobias.funke@reqr...	Tobias	Funke	https://reqres.in...
10	byron.fields@reqr...	Byron	Fields	https://reqres.in...
11	george.edwards@re...	George	Edwards	https://reqres.in...
12	rachel.howell@req...	Rachel	Howell	https://reqres.in...

	Longitude	Latitude	Weather ID	Weather Main	Weather Description	Weather Icon	Base
0	-0.09	51.50	804	Clouds	overcast clouds	04d	stations
1	2.34	48.85	800	Clear	clear sky	01d	stations
2	77.16	28.55	721	Haze	haze	50n	stations
3	72.89	19.13	701	Mist	mist	50n	stations
4	151.20	-33.86	800	Clear	clear sky	01n	stations

Explorer

Search tables

Weather\_Lakehouse

Tables

Fact\_Live\_APIweat...

Files

# How I schedule a refresh in notebook-

To fetch the fresh data daily from a live API using a **Microsoft Fabric notebook**, you can automate the process using a combination of **notebook scheduling**. If you don't schedule a refresh run, it will give you the old data of that day when you manually run the code of notebook.

**Step 1.** In my notebook, I don't want to run each code cell in schedule run because it has checks also like total columns, total rows, show tables. To exclude them from daily run I will freeze those cells.

**Step 2.** Go to 'Run' in top & choose schedule and give details when you want to run the refresh of data.

**Step 3.** I select the schedule as daily once in a day, time of 6:30PM because it is present in UTC which is 12AM in IST and start from 17th Sept till 31st Oct.

**Step 4.** Save the schedule and add, it will show like in image. It will show the next schedule run also.

**Step 5.** Now you will see my data will refresh every day & you can see the recent runs also mentioning manual run or schedule run details.

**Note:** I checked the data by scheduled run in every 7 min & found it is placing the file in Lakehouse every time when it ran. To reduce this, I used to overwrite query to refresh the same file instead of creating new.

Recent runs for JSON\_into\_table

SparkT-SQL

Refresh

FilterColumn Options

Application name	Subm...	Submitt...	Status	Total d...	Run kind...	Livy Id
JSON_into_table_b0e33e36-f2ae-47a1-9267-	09/19/2025, 4:...	2248601@TC...	Success	1 min 10 sec	Scheduled	b0e33e36-f2a...
JSON_into_table_aa8be7e0-d788-4c82-aa18-	09/18/2025, 6:...	2248601@TC...	Failed	23 sec	Scheduled	aa8be7e0-d7...
JSON_into_table_efacebd4-3ec3-4ba7-8b4c-	09/18/2025, 5:...	2248601@TC...	Stopped (sess	47 min 28 sec	Manual	efacebd4-3ec...

## Scheduled run

+ Add schedule

Add schedule

Repeat

Daily

Time of day

06:30 PM

+ Add time

Start date and time

09/17/2025

End date and time

10/31/2025

Time zone

(UTC) Coordinated Universal Time

SaveCancel

## Schedule

Only 20 total schedules can be created and maintained for this item.  
This includes any schedules created using the API that aren't displayed here.

### Refresh status

#### Last successful refresh

09/17/2025, 2:38 PM

## Scheduled run

Every dayOn

Time of day

06:30 PM

Last successful refresh

Not available

Next refresh

8 hour(s) 7 minute(s)

Show more

Time zone: (UTC) Coordinated Universal Time

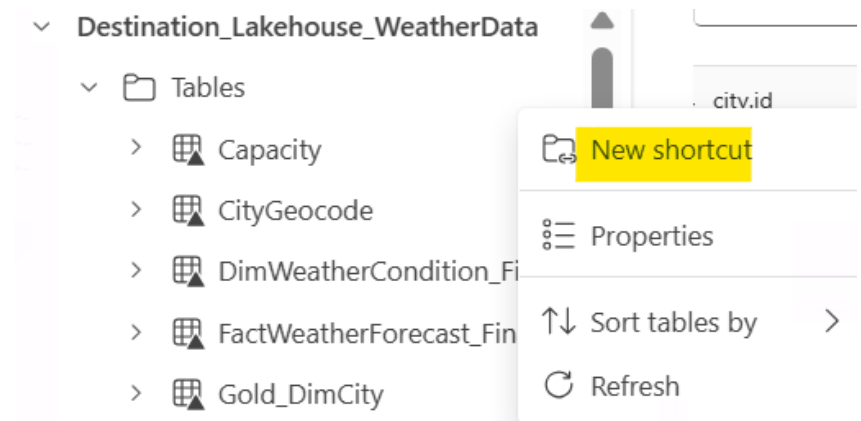
Schedule ID: 748783fe-15d3-4649-8392-7ea37c3388a2

Edit

+ Add schedule

# Shortcut: How I pull data from source Lakehouse & Copy to destination Lakehouse

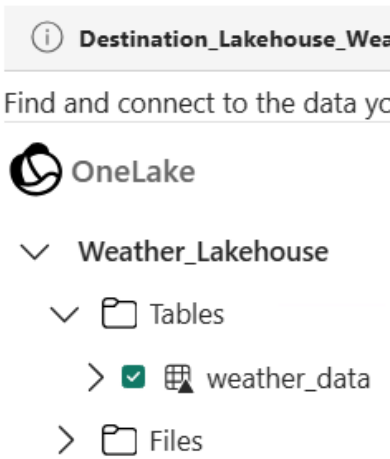
**Step 1.** Inside destination Lakehouse, in tables click on ellipse and create shortcut.



**Step 2.** Choose OneLake & select your source Lakehouse.

**Step 3.** Choose the table/file you want to copy and hit next.

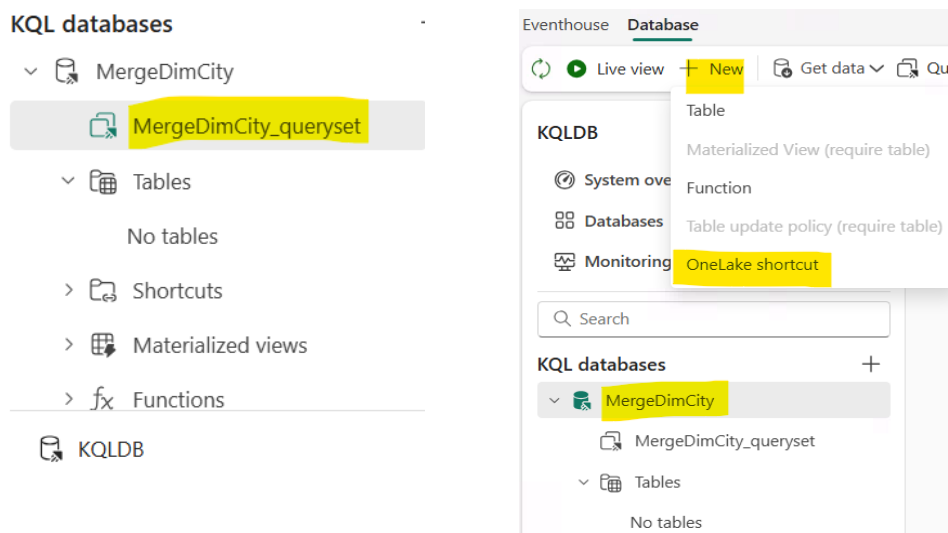
## New shortcut



# How I created KQL Database in MS Fabric

**Event house:** Create or use an existing Event house – this is the container for your KQL databases.

- **KQL Database:** In the Event house Explorer, select **New database** under KQL Databases. Choose between:
  - **New database** (default)
  - **New shortcut database** (follower)
- Once your KQL database is created, an attached **query environment** appears.
- Select the **KQLdatabasename\_queryset** to open the editor.
- Start writing KQL queries to explore your data.



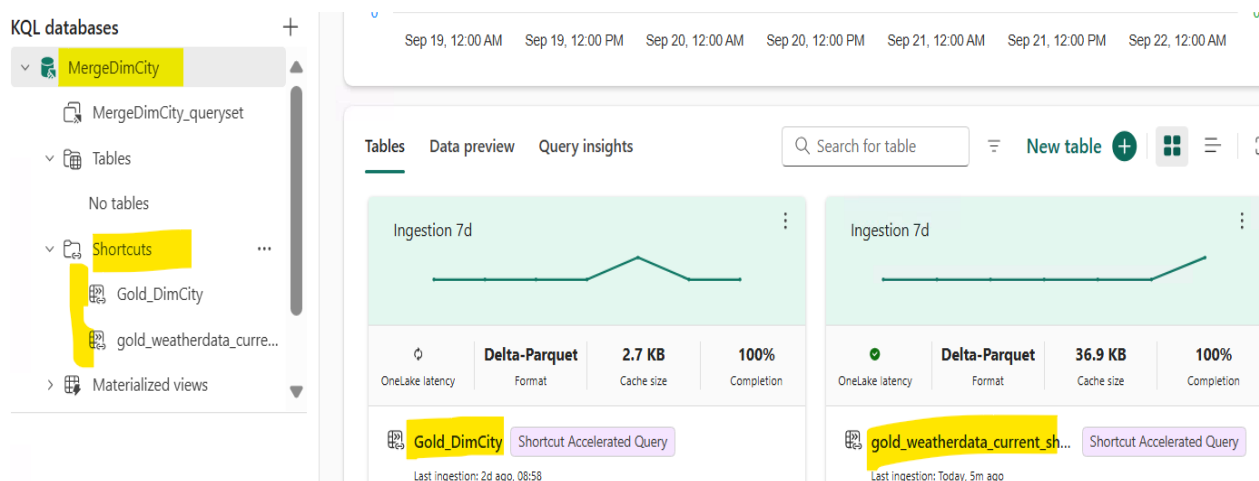
# How to Query Lakehouse Table from KQL Database

**1. Attach the Lakehouse to your KQL Database-** Before querying, ensure your Lakehouse is attached to the KQL database:

- Go to your KQL database in Fabric.
- Click “+ New” → “**OneLake Shortcut**” → Choose OneLake.
- Select the Lakehouse that contains your table & select tables which you want to query in KQL.

This creates a logical link so KQL can access the Delta tables stored in the Lakehouse.

**2. Explore Available Tables -** Once attached, you can preview available tables as shown in image.



# How I Merge the data using KQL in Real-Time Analytics

1. In KQL DB which you created is having in built query set file. Open it to run query in KQL.
2. I firstly gave my table name in query & performed left join operation to get specific columns from another table.

```
external_table('Table A')
| join kind=inner (
    external_table('Table B')
    | project column_names
) on $left.common_column == $right.common_column
| project column_names
```

3. I merged another table which is fact tables of both source Lakehouse & destination Lakehouse using union to add all required columns of both fact data..

```
Union(
    'Table A'
    | project column_names
),
(
    'Table B'
    | project column_names
)
```

KQL databases

▼ MergeDimCity

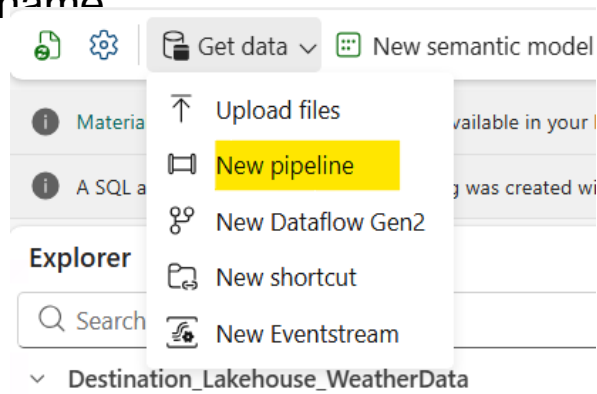
MergeDimCity\_queryset

> Tables

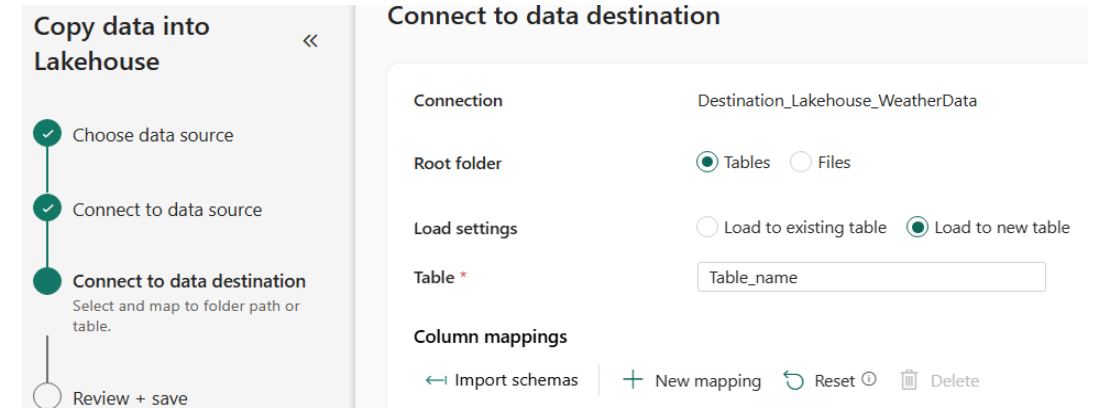


# How I loaded KQL Data Tables in Target Lakehouse using Real-Time Analytics Pipeline

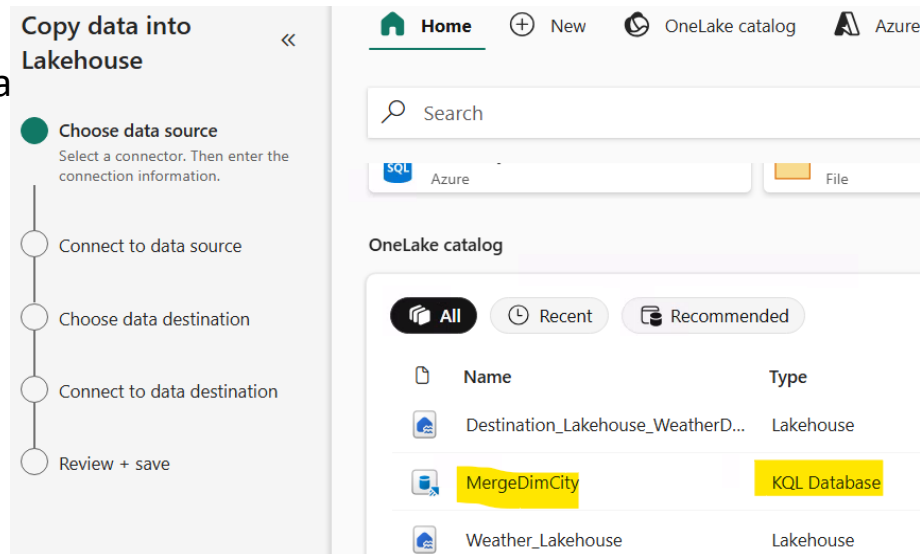
1. In workspace go to destination Lakehouse > Get Data > New Pipeline>Give name



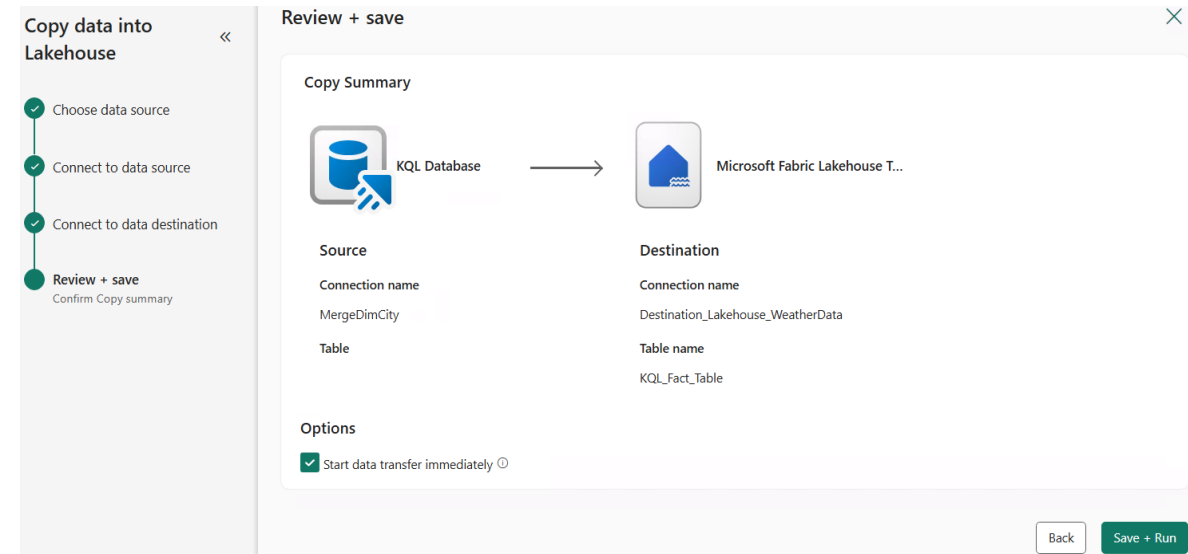
3. Paste the KQL Query & hit next.



2. Select data




4. Select destination lakehouse & give new table name & hit next.

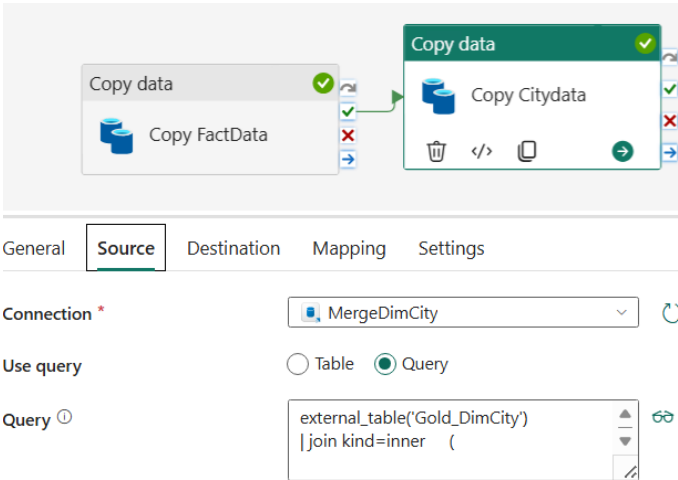


5. I have another query also in KQL which I need to load in table form in destination Lakehouse.

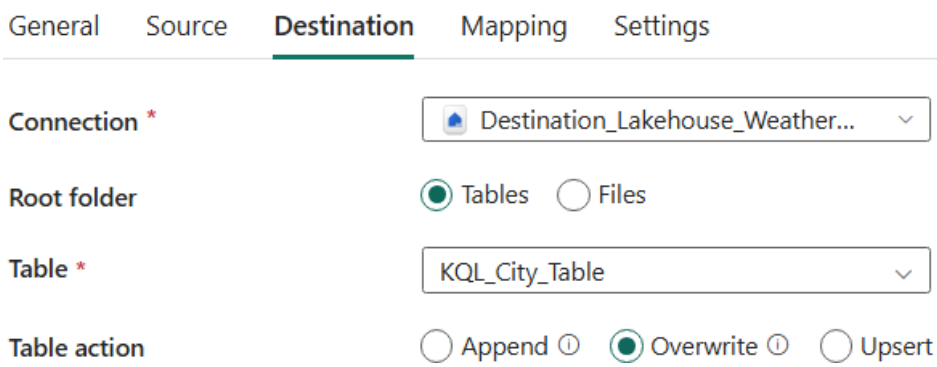
6. Open same pipeline and click on add activity on existing copy data activity block.



7. Click on Copy Data Activity & enter Source as KQL DB & paste that query.

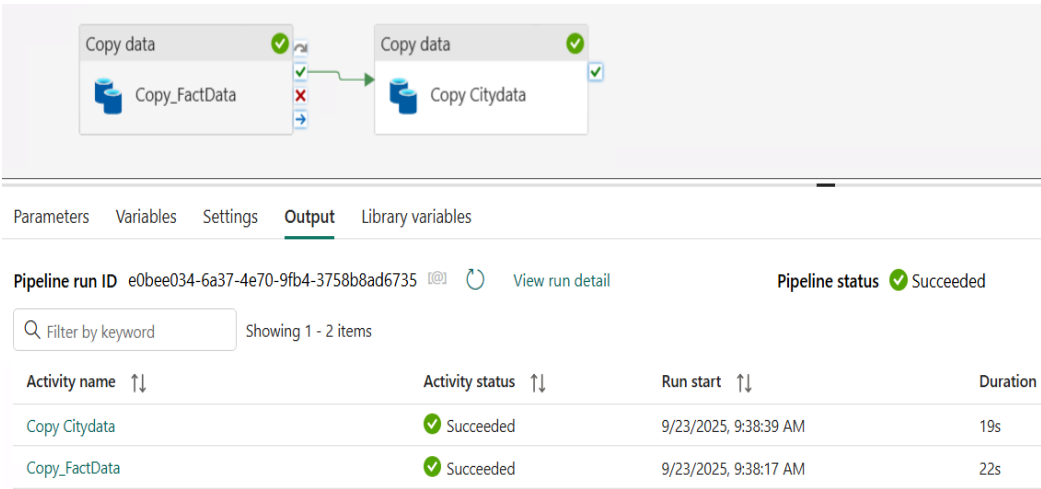


8. Give destination details and give new name in table.







9. Validate & run the pipeline.

10. You can see my tables are loaded in lakehouse after pipeline ran.




# Microsoft Fabric | Ingest Weather API Data










Created Notebooks for data ingestion

WeatherDataIngestion > Notebooks				
	Name	Type	Task	Owner
	API_fileFormat	Notebook	—	Muskan
	JSON_into_table	Notebook	—	Muskan
	Table Details	Notebook	—	Muskan

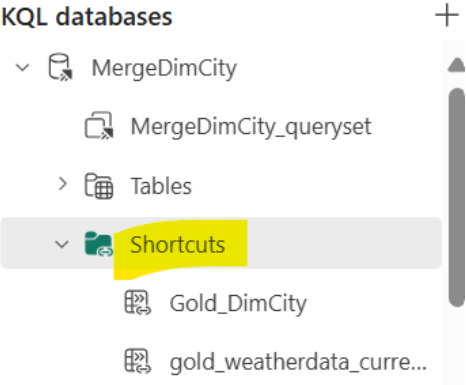
Ingested in Lakehouse using Pyspark

	12 Longitude	12 Latitude	12L WeatherID	ABC WeatherMain	ABC Description	ABC WeatherIcon	ABC Base	12 Temperature
1	-0.09	51.5	804	Clouds	overcast clouds	04n	stations	286.74
2	80.16	13.08	701	Mist	mist	50n	stations	299.51
3	77.6	12.97	803	Clouds	broken clouds	04n	stations	295.19
4	2.34	48.85	803	Clouds	broken clouds	04n	stations	288.65
5	151.2	-33.86	800	Clear	clear sky	01n	stations	284.7

Created shortcut to copy table in silver Lakehouse

Destination_Lakehouse_WeatherData	
∨ 	Tables
> 	Capacity
> 	CityGeocode
> 	DimWeatherCondition_Final
> 	FactWeatherForecast_Final
> 	Gold_DimCity
> 	Gold_DimWeatherCondition
> 	Gold_FactWeatherForecast
> 	gold_weatherdata_current_shortcut

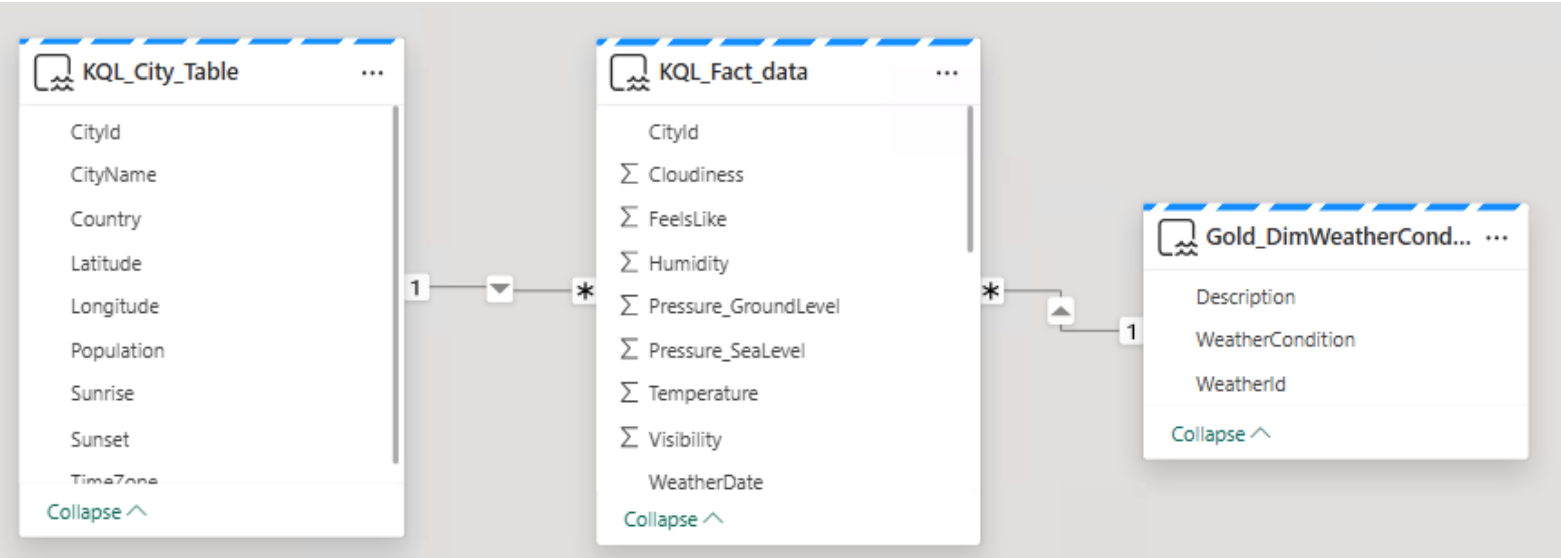
Created shortcut in KQL DB  
to load lakehouse tables



Merged fact data using KQL & load in  
lakehouse

CityId	CityName	Latitude	Longitude	Country	Population	TimeZone	Sunrise	Sunset
6,324,621	Mumbai	19.1316	72.8914	IN	20,000	19,800	2025-09-22 00:57:17	2025-09-22 13:05:13
1,277,333	Bengaluru	12.9705	77.6048	IN	5,104,047	19,800	2025-09-22 00:38:42	2025-09-22 12:46:08
1,278,840	Chennai	13.0861	80.1627	IN	341,049	19,800	2025-09-22 00:28:27	2025-09-22 12:35:54
2,643,741	London	51.5096	-0.0991	GB	7,556,900	3,600	2025-09-21 05:44:55	2025-09-21 18:02:00

Created Semantic Model using Gold  
tables



## KPI for Grid Stability :

KPI	Formula	Thresholds/Insights
Wind Power Potential Index (WPPI)	$\text{wind.speed} / \text{cut\_in\_speed}$ (assume cut-in = 3 m/s)	<1 = no power, 1–3 = good generation
Cloudiness Index (CI)	Avg Cloudiness /100	➤ 0.7 → output drop
Irradiance Proxy (IrrProxy)	$(1 - \text{CI}) \times 1000 \text{ W/m}^2$	Used to forecast solar generation

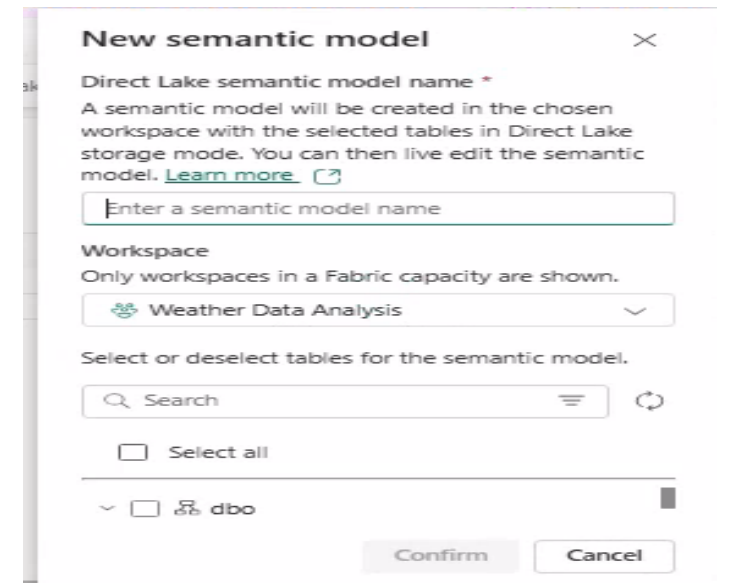
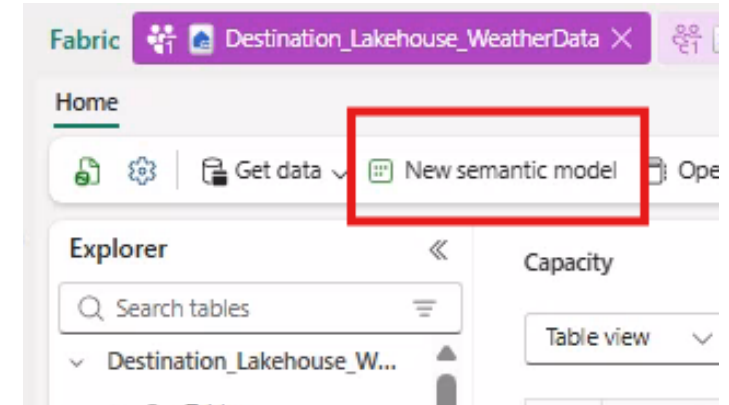
## KPI for Predictive Maintenance:

KPI	Formula	Thresholds/Insights
Temperature Stress Index (TSI)	$(\text{avg temp} - 25) / 25$	High → overheating of inverters/panels
Icing Risk Index (IRI)	If $(-5 \leq \text{temp} \leq 2)$ AND humidity > 80% AND precipitation > 0 then 1 else 0	1 = Icing risk (turbine imbalance/panel loss)
Humidity Stress Index (HSI)	humidity / 100	➤ 0.8 → electrical/corrosion risk
Grid Risk Score	Total number of risky hours/ Total number of hours	0.0 - 0.3 -- low Risk 0.3 - 0.6 -- Medium Risk 0.6 - 1 -- High Risk

# Power BI semantic model :

A **semantic model** in Microsoft Fabric is a logical representation of an analytical domain, designed to facilitate deeper analysis and reporting. It typically follows a star schema, with facts representing the domain and dimensions allowing for detailed analysis, filtering, and calculations.

1. Open the lakehouse and select **New Power BI semantic model** from the ribbon.
2. Enter a name for the new semantic model, select a workspace to save it in, and pick the tables to include. Then select **Confirm**.
3. The new Power BI semantic model can be edited in the workspace, where you can add relationships, measures, rename tables and columns, choose how values are displayed in report visuals, and much more.
4. To edit the Power BI semantic model later, select **Open data model** from the semantic model context menu



# Power BI Report :

